

```
In [3]: from dotenv import load_dotenv
import os

# Load environment variables from .env file
load_dotenv()

NE04J_USER = os.getenv("NE04J_USER")
NE04J_PASSWORD = os.getenv("NE04J_PASS")
openai_api_key = os.getenv("openai_api_key")
```

```
In [ ]: from google.cloud import bigquery
client = bigquery.Client()

admission_query = """
WITH eligible_patients AS (
    SELECT subject_id
    FROM `physionet-data.mimiciii_clinical.admissions`
    GROUP BY subject_id
    HAVING COUNT(DISTINCT hadm_id) BETWEEN 1 AND 3
),
sampled_patients AS (
    SELECT subject_id
    FROM eligible_patients
    ORDER BY RAND()
    LIMIT 1000
)
SELECT a.*
FROM `physionet-data.mimiciii_clinical.admissions` a
JOIN sampled_patients s ON a.subject_id = s.subject_id
WHERE a.hadm_id IS NOT NULL
"""

sampled_patients_admissions_df = client.query(admission_query).to_dataframe()

# Formating the date columns
sampled_patients_admissions_df['ADMITTIME'] = pd.to_datetime(sampled_patient
sampled_patients_admissions_df['DISCHTIME'] = pd.to_datetime(sampled_patient

# Flag readmissions within 30 days
sampled_patients_admissions_df = sampled_patients_admissions_df.sort_values(
sampled_patients_admissions_df["NEXT_ADMITTIME"] = sampled_patients_admissio
sampled_patients_admissions_df["DAYS_TO_NEXT"] = (sampled_patients_admissio
sampled_patients_admissions_df["READMIT_30D"] = sampled_patients_admissions_

sampled_hadm_ids = sampled_patients_admissions_df["HADM_ID"].unique().tolist
sampled_patient_ids = sampled_patients_admissions_df['SUBJECT_ID'].unique().

# Convert list of Subject and HADM IDs into a SQL IN clause
id_list_str = ", ".join(str(x) for x in sampled_patient_ids)
id_filter = f"({id_list_str})"
```

```

hadm_list_str = ", ".join(str(x) for x in sampled_hadm_ids)
hadm_filter = f"({hadm_list_str})"

```

```

In [106... patient_query = """
SELECT *
FROM `physionet-data.mimiciii_clinical.patients` p
WHERE p.subject_id IN {id_filter}
"""

# Download the sampled patients data to CSV
sampled_patients_df = client.query(patient_query.format(id_filter=id_filter))
sampled_patients_df.to_csv("./neo4j-community-2025.03.0/import/patients.csv")

# Add AGE column to admissions dataframe
sampled_patients_admissions_df = sampled_patients_admissions_df.merge(sample

# Compute age at admission; Account for MIMIC III obfuscation of ages over 8
sampled_patients_admissions_df["AGE"] = (sampled_patients_admissions_df["ADM
sampled_patients_admissions_df["AGE"] = sampled_patients_admissions_df["AGE"

# Drop DOB column after use
sampled_patients_admissions_df = sampled_patients_admissions_df.drop(columns

# Download the sampled admissions data to CSV
sampled_patients_admissions_df.to_csv("./neo4j-community-2025.03.0/import/ac

```

/Users/abemankavil/Desktop/MS in AI/AI395T-AIinHealthcare/HighRisk_Project/HighRiskEnv/lib/python3.9/site-packages/google/cloud/bigquery/table.py:1933: UserWarning: BigQuery Storage module not found, fetch data with the REST endpoint instead.

warnings.warn(

```

In [107... diag_query = """
SELECT d.subject_id, d.hadm_id, d.icd9_code, icd.short_title AS diagnosis
FROM `physionet-data.mimiciii_clinical.diagnoses_icd` d
JOIN `physionet-data.mimiciii_clinical.d_icd_diagnoses` icd
    ON d.icd9_code = icd.icd9_code
WHERE d.hadm_id IN {hadm_filter}
"""

diag_df = client.query(diag_query.format(hadm_filter=hadm_filter)).to_dataframe

diag_df.to_csv("./neo4j-community-2025.03.0/import/diagnoses.csv", index=False)

```

/Users/abemankavil/Desktop/MS in AI/AI395T-AIinHealthcare/HighRisk_Project/HighRiskEnv/lib/python3.9/site-packages/google/cloud/bigquery/table.py:1933: UserWarning: BigQuery Storage module not found, fetch data with the REST endpoint instead.

warnings.warn(

```

In [108... drug_query = """
SELECT subject_id, hadm_id, drug, route, startdate, enddate, dose_val_rx, dose_unit_rx
FROM `physionet-data.mimiciii_clinical.prescriptions`
WHERE hadm_id IN {hadm_filter}
AND drug IS NOT NULL
"""

```

```
drug_df = client.query(drug_query.format(hadm_filter=hadm_filter)).to_dataframe()
drug_df.to_csv("./neo4j-community-2025.03.0/import/prescriptions.csv", index
```

/Users/abemankavil/Desktop/MS in AI/AI395T-AIinHealthcare/HighRisk_Project/HighRiskEnv/lib/python3.9/site-packages/google/cloud/bigquery/table.py:1933: UserWarning: BigQuery Storage module not found, fetch data with the REST endpoint instead.
warnings.warn(

```
In [109... lab_query = """
SELECT le.subject_id, le.hadm_id, le.charttime, d.label AS lab_name, le.value AS lab_value
FROM `physionet-data.mimiciii_clinical.labevents` le
JOIN `physionet-data.mimiciii_clinical.d_labitems` d
  ON le.itemid = d.itemid
WHERE le.hadm_id IN {hadm_filter}
AND le.valuenum IS NOT NULL

"""

lab_df = client.query(lab_query.format(hadm_filter=hadm_filter)).to_dataframe()
lab_df.to_csv("./neo4j-community-2025.03.0/import/labevents.csv", index=False)
```

/Users/abemankavil/Desktop/MS in AI/AI395T-AIinHealthcare/HighRisk_Project/HighRiskEnv/lib/python3.9/site-packages/google/cloud/bigquery/table.py:1933: UserWarning: BigQuery Storage module not found, fetch data with the REST endpoint instead.
warnings.warn(

```
In [1]: from neo4j import GraphDatabase
from langchain.graphs import Neo4jGraph
from langchain.chat_models import ChatOpenAI
from langchain.chains import GraphCypherQAChain
from sklearn.ensemble import IsolationForest
import numpy as np
from sklearn.manifold import TSNE
import matplotlib.pyplot as plt
import pandas as pd
from openai import OpenAI
from graphdatascience import GraphDataScience
import seaborn as sns
```

/Users/abemankavil/Desktop/MS in AI/AI395T-AIinHealthcare/EHR_Anomaly_Detection/HighRiskEnv/lib/python3.9/site-packages/tqdm/auto.py:21: TqdmWarning: IPywidgets not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm

```
In [ ]: # -----
# Step 1: Neo4j + OpenAI Setup
# -----

NEO4J_URL = "bolt://localhost:7687"

# LangChain graph wrapper
graph_langchain = Neo4jGraph(
```

```

url=NEO4J_URL,
username=NEO4J_USER,
password=NEO4J_PASSWORD
)

# Neo4j driver and GDS interface
gds = GraphDataScience(NEO4J_URL, auth=(NEO4J_USER, NEO4J_PASSWORD))

```

/var/folders/0r/lprjchvd593dyrpzsr6jjrm0000gn/T/ipykernel_83891/3600210996.py:13: LangChainDeprecationWarning: The class `Neo4jGraph` was deprecated in LangChain 0.3.8 and will be removed in 1.0. An updated version of the class exists in the :class:`~langchain-neo4j` package and should be used instead. To use it run `pip install -U :class:`~langchain-neo4j` and import as `from :class:`~langchain_neo4j import Neo4jGraph`.

```
graph_langchain = Neo4jGraph(
```

```

In [4]: # -----
# Step 2: Generate Embeddings Using GDS Python Client
# -----

def generate_node2vec_embeddings():
    if gds.graph.exists("admission_subgraph").get("exists", False):
        gds.graph.drop("admission_subgraph")

    G, result = gds.graph.project(
        "admission_subgraph",
        node_spec={"Admission": {},
                  "Diagnosis": {},
                  "LabResult": {},
                  "Medication": {}},
        relationship_spec={
            "HAS_DIAGNOSIS": {"orientation": "UNDIRECTED"},
            "HAS_LAB": {"orientation": "UNDIRECTED"},
            "HAS_MEDICATION": {"orientation": "UNDIRECTED"}
        }
    )

    print(f"Projection created in {result['projectMillis']} ms")

    result = gds.node2vec.write(
        G, # Graph object from projection
        writeProperty="embedding",
        embeddingDimension=64,
        iterations=10,
        nodeLabels=["Admission"],
        concurrency=4
    )
    print(f"Node2Vec computed for {result['nodeCount']} nodes")

```

```

In [5]: # -----
# Step 3: Fetch Embeddings from Neo4j
# -----

def get_admission_features():
    query = """
    MATCH (p:Patient)-[:HAS_ADMISSION]->(a:Admission)

```

```

OPTIONAL MATCH (a)-[:HAS_MEDICATION]->(m:Medication)
OPTIONAL MATCH (a)-[:HAS_LAB]->(l:LabResult)
OPTIONAL MATCH (a)-[:HAS_DIAGNOSIS]->(d:Diagnosis)
WHERE a.embedding IS NOT NULL
RETURN a.HADM_ID AS hadm_id, a.embedding AS embedding, a.admission_type
      p.GENDER AS gender, a.ethnicity AS ethnicity, a.age AS age, a.hos
      collect(DISTINCT {name: m.DRUG, dose: m.dose, unit: m.unit}) AS m
      """"

result = gds.run_cypher(query)
df = pd.DataFrame(result)
df = df[df["embedding"].notnull()].copy()
df["embedding"] = df["embedding"].apply(lambda x: [float(i) for i in x])
return df

```

```

In [98]: # -----
# Step 4: Detect Anomalies Using Isolation Forest
# -----
from sklearn.preprocessing import OneHotEncoder, MultiLabelBinarizer

def detect_anomalies(df):
    cat_features = df[["gender", "ethnicity", "admission_type", "hospital_ex
    enc = OneHotEncoder(sparse_output=False, handle_unknown="ignore")
    cat_encoded = enc.fit_transform(cat_features)

    meds_name_enc = MultiLabelBinarizer()
    meds_dose_enc = MultiLabelBinarizer()
    labs_name_enc = MultiLabelBinarizer()
    labs_val_enc = MultiLabelBinarizer()
    diag_enc = MultiLabelBinarizer()

    # Split medications into name and (dose+unit)
    def split_medications(meds):
        names = [m["name"].lower() for m in meds if m.get("name")]
        doses = [f"{m.get('dose', '?')} {m.get('unit', '?')}" for m in meds
        return pd.Series([names, doses])

    # Split labs into name and (value+unit)
    def split_labs(labs):
        names = [l["name"].lower() for l in labs if l.get("name")]
        values = [f"{l.get('value', '?')} {l.get('unit', '?')}" for l in labs
        return pd.Series([names, values])

    df[["med_names", "med_doses"]] = df["meds"].apply(split_medications)
    df[["lab_names", "lab_values"]] = df["labs"].apply(split_labs)

    med_name_feats = meds_name_enc.fit_transform(df["med_names"])
    med_dose_feats = meds_dose_enc.fit_transform(df["med_doses"])
    lab_name_feats = labs_name_enc.fit_transform(df["lab_names"])
    lab_val_feats = labs_val_enc.fit_transform(df["lab_values"])
    diag_feats = diag_enc.fit_transform(df["diag_codes"])

    age_vals = df["age"].fillna(0).values.reshape(-1, 1)
    embed_vals = np.vstack(df["embedding"].values)

```

```

X = np.hstack([embed_vals, age_vals, cat_encoded, med_dose_feats, med_na

model = IsolationForest(n_estimators=100, contamination=0.05, random_sta
df['anomaly_score'] = model.fit_predict(X)
df['anomaly'] = df['anomaly_score'] == -1
df['score'] = model.decision_function(X)

X_tsne = TSNE(n_components=2, perplexity=30, learning_rate=200, n_iter=1

# Plot
plt.figure(figsize=(10, 6))
scatter = plt.scatter(
    X_tsne[:, 0], X_tsne[:, 1],
    c=df["anomaly"],
    cmap="coolwarm",
    alpha=0.7,
    edgecolor='k',
    linewidth=0.3
)

plt.title("t-SNE Projection of EHR Admissions (Anomalies Highlighted)")
plt.xlabel("t-SNE Dimension 1")
plt.ylabel("t-SNE Dimension 2")
plt.legend(*scatter.legend_elements(), title="Anomaly = 1", loc="lower r
plt.grid(True, linestyle="--", alpha=0.3)
plt.tight_layout()
plt.show()

plt.figure(figsize=(10, 6))
sns.histplot(df["score"], bins=50, kde=True)
plt.axvline(df[df['anomaly'] == True]['score'].mean(), color='red', line
plt.title("Distribution of Anomaly Scores (Isolation Forest)")
plt.xlabel("Anomaly Score")
plt.ylabel("Frequency")
plt.legend()
plt.tight_layout()
plt.show()

return df[df['anomaly'] == True].sort_values("score")

```

```

In [ ]: # -----
# Step 5: LLM-Based Explanation via LangChain
# -----

llm = ChatOpenAI(
    model="gpt-4o-mini",
    temperature=0.3,
    openai_api_key=openai_api_key
)

# GraphCypherQChain Setup
qa_chain = GraphCypherQChain.from_llm(
    llm=llm,
    graph=graph_langchain,

```

```

        verbose=False,
        allow_dangerous_requests=True
    )

```

```

def explain_anomaly(hadm_id):
    question = f"""

```

You are a medical graph reasoning assistant. A patient admission is being ev

Only use the following information in your Cypher query:

- Admission attributes: admission_type, diagnosis_free_text, age, ethnicity,
- Patient attributes: GENDER
- Diagnosis: CODE
- LabResult: TEST_NAME, value, unit
- Medication: DRUG, dose, unit

Ensure you **aggregate and deduplicate** all relevant values using `COLLECT`

- Medications: `collect(DISTINCT {{name: m.DRUG, dose: m.dose, unit: m.unit}}`
- Lab Results: `collect(DISTINCT {{name: l.TEST_NAME, value: l.value, unit: l.unit}}`
- Diagnosis: `collect(DISTINCT d.CODE) AS diag_codes`

Explain why the hospital admission with HADM_ID {hadm_id} might be clinically
Identify 2 to 3 clinically meaningful anomalies for this hospital admission.

Your analysis can consider, but is not limited to, any of the following:

- Unusual medication–diagnosis combinations
- Rare or unexpected lab test patterns based on admission type
- Mismatches between patient demographics (e.g., age, gender, ethnicity) and
- Co-occurrences of diagnoses that are atypical when found together

You may also identify other patterns or inconsistencies you find clinically

"""

```

    try:
        return qa_chain.invoke(question)['result']
    except Exception as e:
        return f"LLM explanation failed: {e}"

```

```

In [8]: # -----
# Step 6: Run Full Pipeline
# -----

```

```

print("[1] Generating Node2Vec embeddings using GDS Python client...")
generate_node2vec_embeddings()

```

```

[1] Generating Node2Vec embeddings using GDS Python client...
Projection created in 816 ms
Node2Vec computed for 1163 nodes

```

```

In [9]: print("[2] Fetching embeddings + node properties from Neo4j...")
df = get_admission_features()

```

```

[2] Fetching embeddings + node properties from Neo4j...

```

```
/Users/abemankavil/Desktop/MS in AI/AI395T-AIinHealthcare/HighRisk_Project/HighRiskEnv/lib/python3.9/site-packages/graphdatascience/query_runner/neo4j_query_runner.py:277: RuntimeWarning: null value eliminated in set function.  
warnings.warn(warning)
```

```
In [77]: df.head()
```

Out[77]:

	hadm_id	embedding	admission_type	diagnosis	gender
0	144471	[-0.0001035963068716228, -0.007562919519841671...	EMERGENCY	WEAKNESS	F
1	162937	[-0.0008790345746092498, 0.0069307610392570496...	EMERGENCY	MULTI MYELOMA;PAIN	F
2	169482	[-0.0008776350878179073, 0.004281629808247089,...	EMERGENCY	ACUTE GASTROENTERITIS	F
3	177186	[0.0010735585819929838, -0.0011072519700974226...	EMERGENCY	ALTERED MENTAL STATUS	F
4	127794	[0.0007670226041227579, 0.0009274794138036668,...	ELECTIVE	ANEURYSM/SDA	F

5 rows x 21 columns


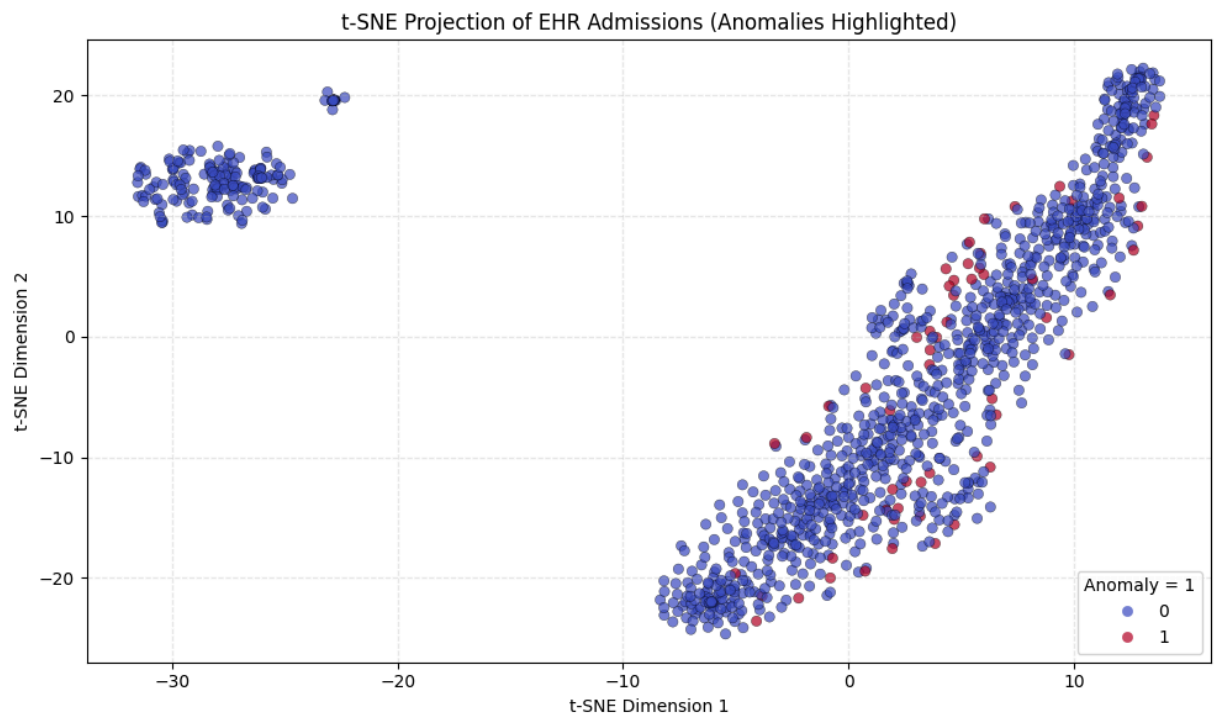


```
In [99]: print("[3] Running Isolation Forest to detect anomalies...")  
anomalies = detect_anomalies(df)  
print(f"Detected {len(anomalies)} anomalous admissions.\n")
```

[3] Running Isolation Forest to detect anomalies...

/Users/abemankavil/Desktop/MS in AI/AI395T-AIinHealthcare/HighRisk_Project/HighRiskEnv/lib/python3.9/site-packages/sklearn/manifold/_t_sne.py:1164: FutureWarning: 'n_iter' was renamed to 'max_iter' in version 1.5 and will be removed in 1.7.

warnings.warn(



Detected 59 anomalous admissions.

```
In [25]: import concurrent.futures

def explain_anomaly_wrapper(hadm_id):
    try:
        explanation = explain_anomaly(hadm_id)
```

```
        return hadm_id, explanation
    except Exception as e:
        return hadm_id, f"Error: {e}"

N = int(input("How many top-N anomalies would you like to explain? "))
top_anomalies = anomalies.head(N)

print("[4] Generating explanations for top anomalies...\n")

with concurrent.futures.ThreadPoolExecutor() as executor:
    futures = {executor.submit(explain_anomaly_wrapper, row['hadm_id']): row
               for row in top_anomalies}
    for future in concurrent.futures.as_completed(futures):
        hadm_id, explanation = future.result(timeout=90)
        print(f"HADM_ID: {hadm_id}")
        print("Explanation:", explanation)
        print("-" * 60)
```

[4] Generating explanations for top anomalies...

HADM_ID: 164694

Explanation: The hospital admission with HADM_ID 164694 may be clinically anomalous for several reasons:

1. ****Diagnosis and Age Mismatch****: The patient is a 38-year-old male diagnosed with Acute Myelogenous Leukemia (AML), which is more commonly diagnosed in older adults. This diagnosis at a relatively young age could be considered unusual and warrants further investigation into potential underlying genetic or environmental factors.
2. ****Medication-Diagnosis Combination****: The patient is receiving a wide array of medications, including Quetiapine Fumarate, which is primarily used for psychiatric conditions. This raises concerns about the appropriateness of this medication in the context of AML treatment, as it may not be a standard part of the therapeutic regimen for this diagnosis.
3. ****Lab Test Patterns****: The lab results show a notably low platelet count (29.0 K/uL), which is expected in patients with leukemia. However, the presence of atypical lymphocytes (10.0%) and a high lactate dehydrogenase (LD) level (1168.0 IU/L) could indicate a more severe underlying condition or complications, such as tumor lysis syndrome or an infection, which are concerning in the context of AML.

These anomalies suggest that the patient's clinical presentation may not align with typical expectations for their diagnosis and demographic profile, indicating a need for closer scrutiny and possibly a reevaluation of their treatment plan.

HADM_ID: 119862

Explanation: The hospital admission with HADM_ID 119862 may be clinically anomalous for several reasons:

1. ****Unusual Medication-Diagnosis Combinations****: The patient, a 57-year-old Black/Haitian male admitted for fever under emergency conditions, received a wide array of medications that may not typically be associated with a straightforward fever diagnosis. For instance, the administration of medications such as Thrombin and Filgrastim suggests a potential underlying condition that requires hemostatic support or stimulation of bone marrow activity, which is not common for a fever alone.
2. ****Rare Lab Test Patterns****: The lab results indicate significant abnormalities, such as a high level of Alanine Aminotransferase (ALT) at 940 IU/L, which is markedly elevated and may suggest liver dysfunction or damage. Additionally, the presence of elevated lactate levels (up to 12.0 mmol/L) could indicate tissue hypoperfusion or sepsis, which aligns with the emergency admission but raises concerns about the patient's overall stability and potential underlying conditions.
3. ****Mismatches Between Patient Demographics and Clinical Presentation****: The patient's age and ethnicity may also present a clinical anomaly when considering the typical causes of fever in this demographic. For instance, certain infections or conditions that commonly cause fever may be less prevalent or present atypically in older Black/Haitian males, suggesting a need for further investigation into the underlying causes of the fever.

These anomalies warrant further clinical evaluation to ensure appropriate diagnosis and treatment, as they may indicate more complex underlying health issues.

HADM_ID: 131328

Explanation: The hospital admission with HADM_ID 131328 may be clinically anomalous for several reasons:

1. ****Unusual Medication-Diagnosis Combinations****: The patient, a 53-year-old white female diagnosed with acute myeloid leukemia and undergoing a bone marrow transplant, is prescribed a wide array of medications, including high doses of chemotherapeutic agents like Cyclophosphamide (4860 mg) and Mitoxantrone (15 mg). Such high dosages, particularly in combination with other medications like Fentanyl Citrate (2.5 mg) and Morphine Sulfate (15 mg), raise concerns about potential drug interactions and the management of pain and side effects during a critical treatment phase.
2. ****Rare Lab Test Patterns****: The lab results show a significant variation in key indicators. For instance, the patient's creatinine levels fluctuate between 0.4 mg/dL and 2.0 mg/dL, indicating possible acute kidney injury or renal impairment, which is not uncommon in patients undergoing chemotherapy. Additionally, the presence of atypical lymphocytes (0.0% to 1.0%) alongside elevated levels of lactate dehydrogenase (up to 1290 IU/L) may suggest underlying complications or an unusual response to treatment.
3. ****Demographic Mismatches****: The patient's age and gender may also present anomalies. While acute myeloid leukemia can occur in older adults, the combination of a relatively young female patient with such a severe diagnosis and the aggressive treatment protocol raises questions about the appropriateness of the treatment plan, especially considering the potential for adverse effects in this demographic.

These factors collectively suggest that the admission may involve complexities that warrant further investigation to ensure optimal patient care and management.
